

A Model-Based Algorithm for Optimizing I/O Intensive Applications in Clouds using VM-Based Migration

Kento Sato† Hitoshi Sato† Satoshi Matsuoka†‡

†: Tokyo Institute of Technology


‡: National Institute of Informatics



Outline

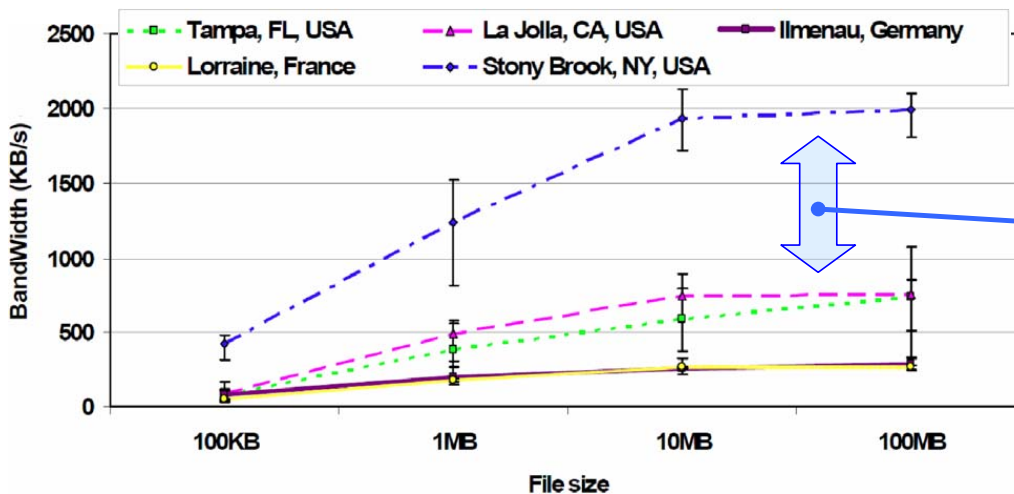
- Introduction
- Target cloud model
- Proposal
 - DAG algorithm
 - Markov model
 - Performance model
- Evaluation
- Conclusion

Background

- Large-scale distributed file system
 - Providing much larger amounts of storage resources than those of typical single-site
 - Giving a common view of all files stored independent from which node access the data
- Amazon S3(simple storage service) 
 - virtually infinite storage spaces with high availability
 - cost-effective pay-as-you-go model

What's the problem ?

- Data Transfer Cost
 - Causing I/O performance degradation of data intensive applications



the location of the client impacts on the observed data access performance in Amazon S3



- Previous approach: File migration
 - File replication & File caching

Graph: Quoted from [Amazon S3 for Science Grids: a Viable Solution ?](#)
in 4th USENIX Symposium on Networked Systems Design & Implementation (NSDI'07), 2007.
[M. Palankar, A. Onibokun, et al.](#)

Our Approach

- VM-Based Approach: VM migration
 - Being in practical use
 - Migrating VMs onto the locations that hold target files
 - 😊 Increasing the performance of file accesses
 - 😞 Causing also VM migration cost
- ⇒ Difficult to determine when and where to migration VMs

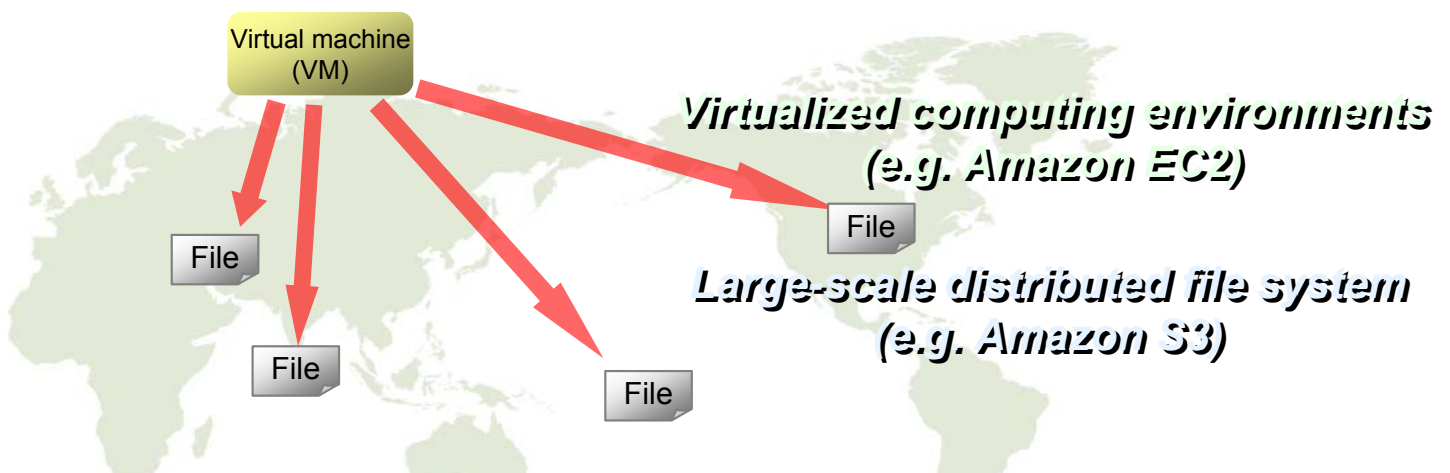
Represent VM's file access patterns as a DAG,
and determine the best location for file access

Goal and Achievement

- Goal
 - Optimization of I/O intensive application in Cloud using VM-based migration
- Achievement
 - Proposed a model-based I/O performance optimization algorithm for data-intensive application
 - Our algorithm can achieve higher I/O performance than simple techniques
 - Never migrating VM: 38%
 - Always migrating VM: 47%

Our Target Cloud Model

- Virtualized computing environments on distributed file system
- Target jobs feature: data-intensive application that accesses distributed multi-files
 - write-once, read-mostly applications



Optimizing the jobs by improving read performance

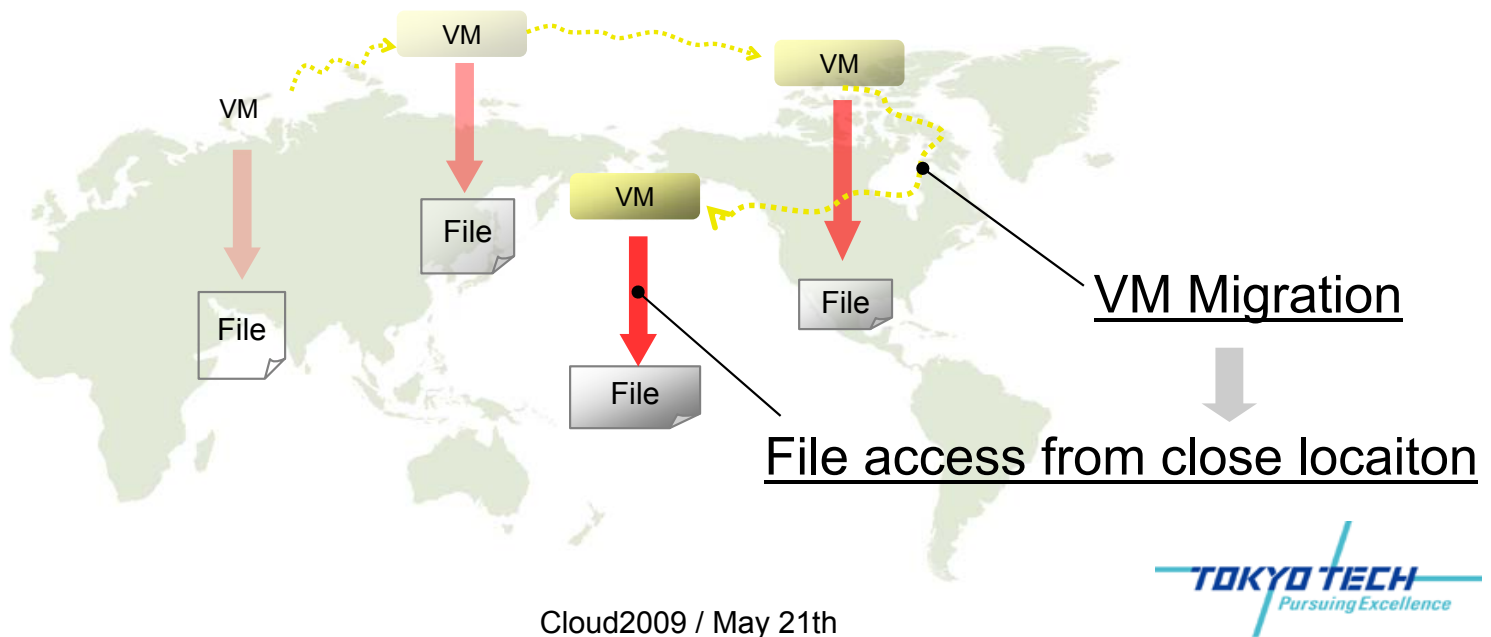
Previous Approach

- **File replication & caching** [Venugopal et al. '06]
 - 😊 Minimizing remote file accesses by creating multiple copies and caching frequently-accessed hot file
 - 😞 Introducing a large amount of file transfer and storage consumption
- **File-location-aware job scheduling** [Shankar et al. '07]
 - 😊 Submitting jobs to sites where target files are located to avoid remote file access
 - 😞 Still causing remote file access, in case a job accesses to geographically distributed files

Our Approach

Migrates VM to onto close locations to target files

- Expected to improve the I/O performance




Difficulty of VM relocation algorithm

Considering the overhead of VM migration

- Not good to migrate VM to target files every times
- File access time and VM migration time depends on runtime environments
 - e.g.) Network throughputs, access file size, VM memory size etc

⇒ We have to determine the optimal migration strategy from the runtime environments




Optimal VM relocation techniques

- Determine VM migration strategies
 - i.e. When VM should be migrated to which sites
 - Minimizing file access time including VM migration time
- Collection of Information to be used
 - Cloud Information:
 - inter-site throughputs, local file system throughputs within each site
 - File Information:
 - size, location, dependency
 - VM Information:
 - memory size, location
- Output a optimal location for requested file


Overview of our algorithm

- Representing possible VM location as a DAG
 - Vertex: File access location
 - Edge: VM migration
- Calculating shortest path of the DAG
 - Vertex weights: Expected file access time
 - Edge weights: VM migration time





Example

- Consider a simple situation
 - File location: File 1 (Site B), File 2 (Site C)
 - VM location: site A
- Explain how to determine a optimal location of VM that access to File 1



Possible migration strategies

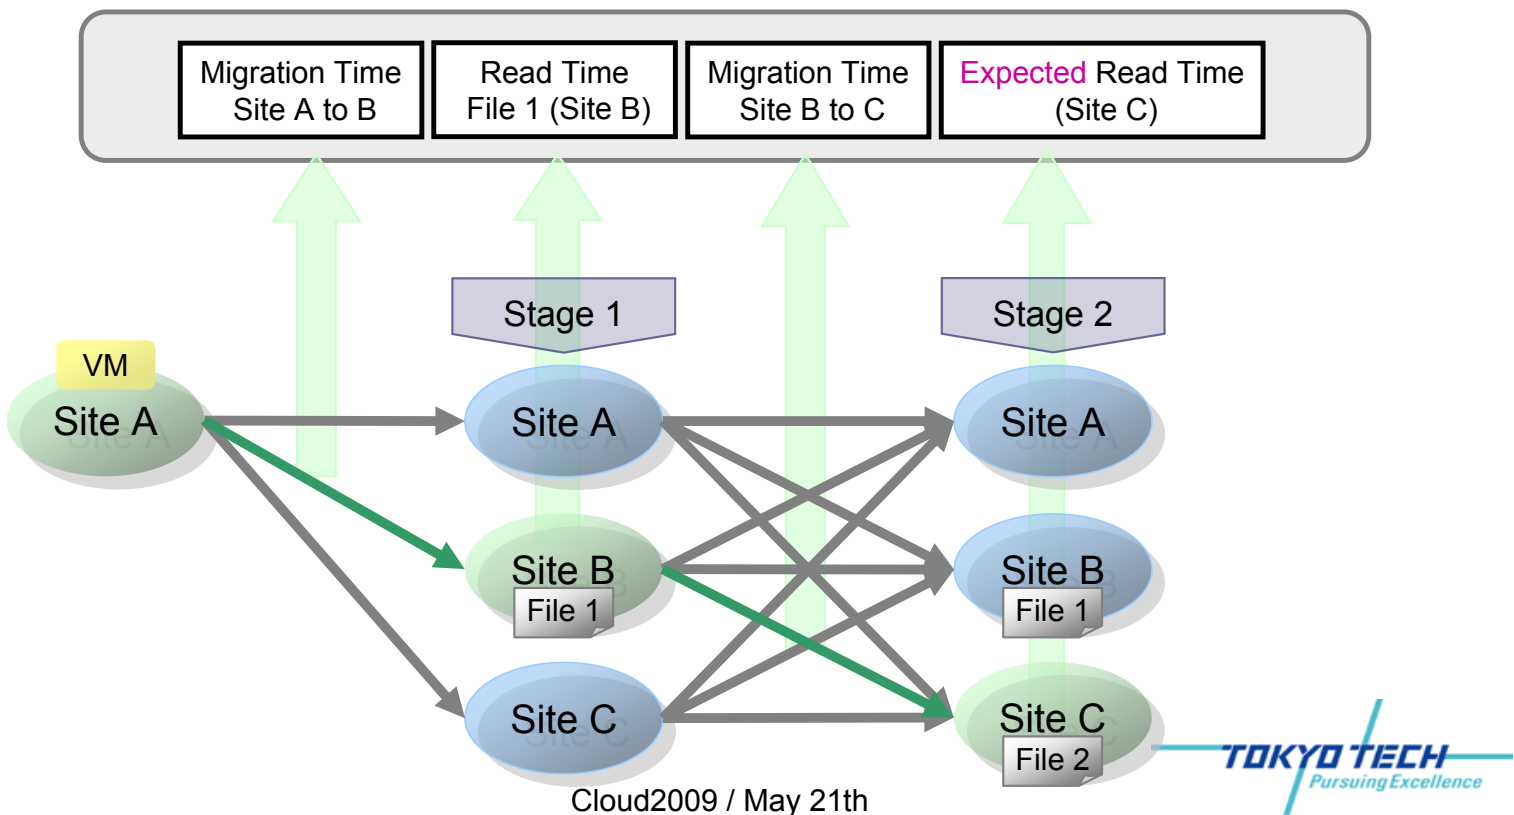
- Representing possible VM location as a DAG
 - Vertex: File access location
 - Edge: VM migration




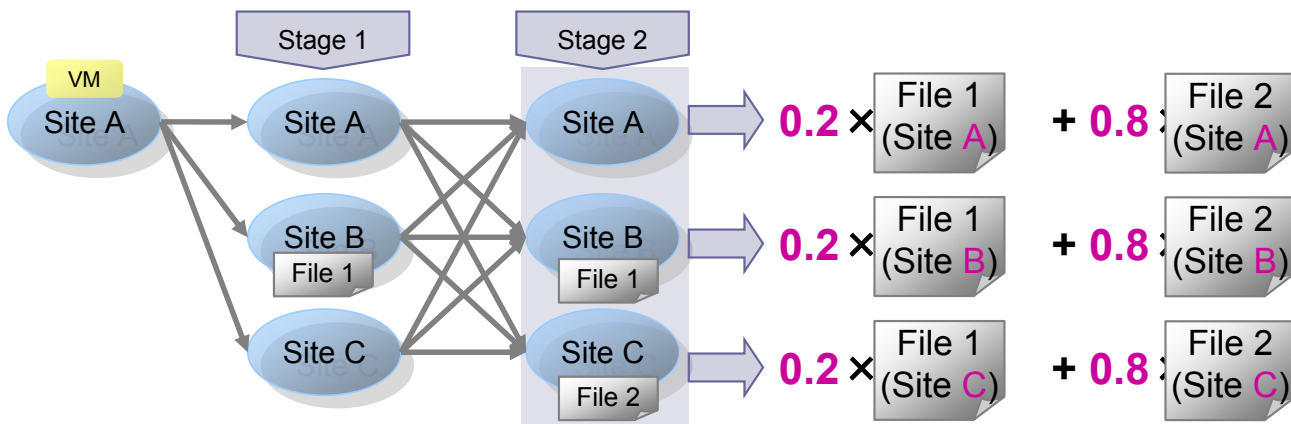
Cloud2009 / May 21th

How to calculate the total access time ?

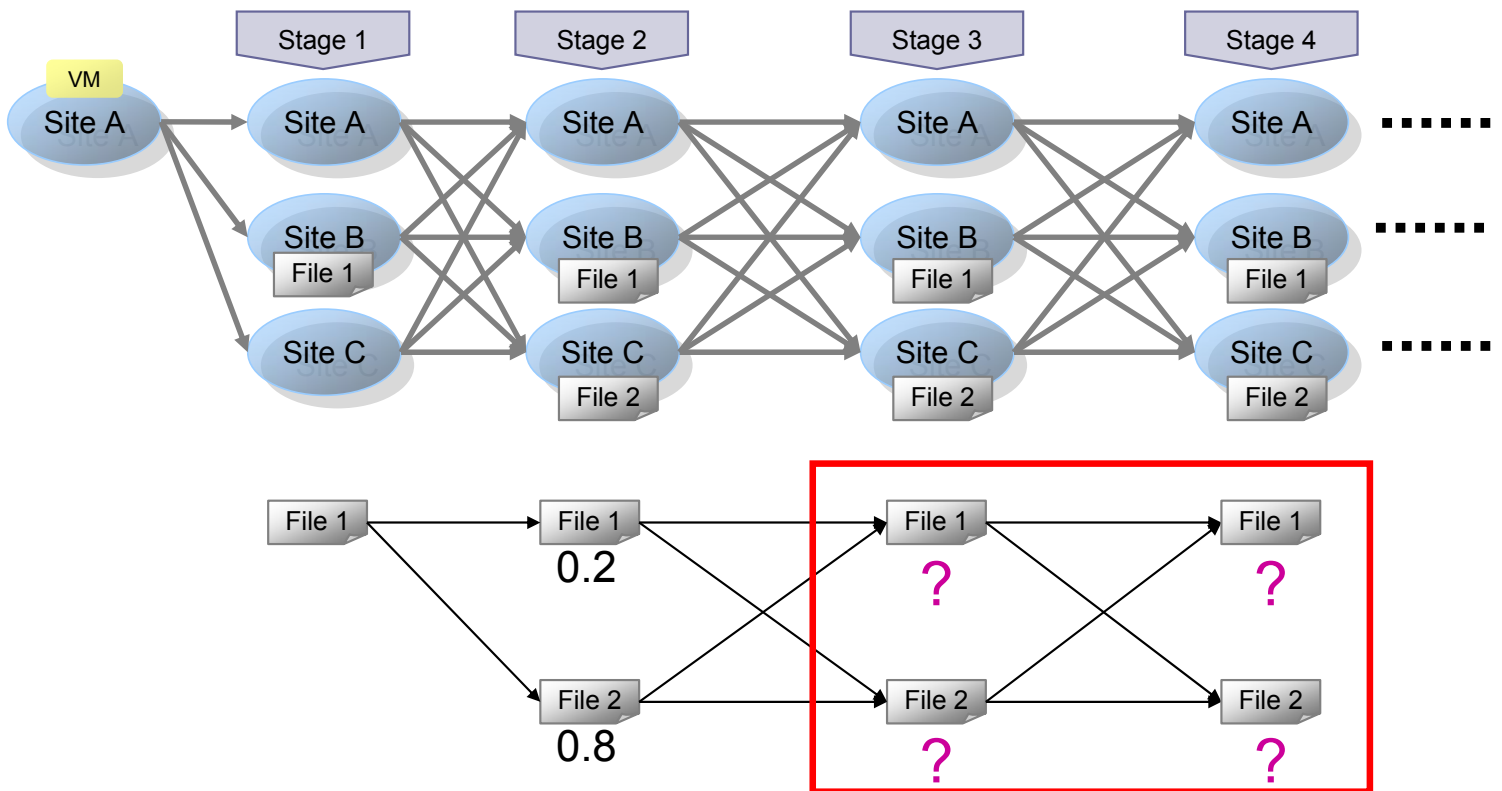
The total access time is the summation of following times



How to calculate the expected file access time ?



How to calculate expected file access time on the other stages ?




File Access Markov Model

Calculating expected file access time from markov model


- Markov model

- representing the probability of access transitions from one file to another from monitored trace




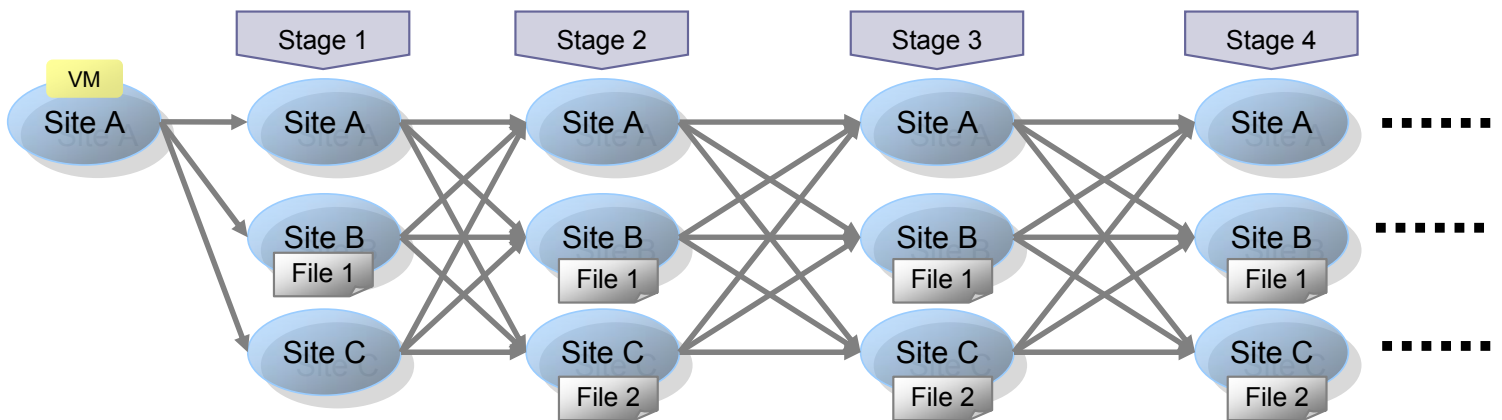
- Stochastic matrix

- Describing a markov model as a matrix
- P^k : the possibility of file access transitoins from one file to another with k -step


$$P = \begin{matrix} & \begin{matrix} \text{File 1} & \text{File 2} \end{matrix} \\ \begin{matrix} \text{File 1} \\ \text{File 2} \end{matrix} & \begin{pmatrix} 0.2 & 0.8 \\ 1.0 & 0.0 \end{pmatrix} \end{matrix}$$

How to calculate expected file access time on the other stages ?

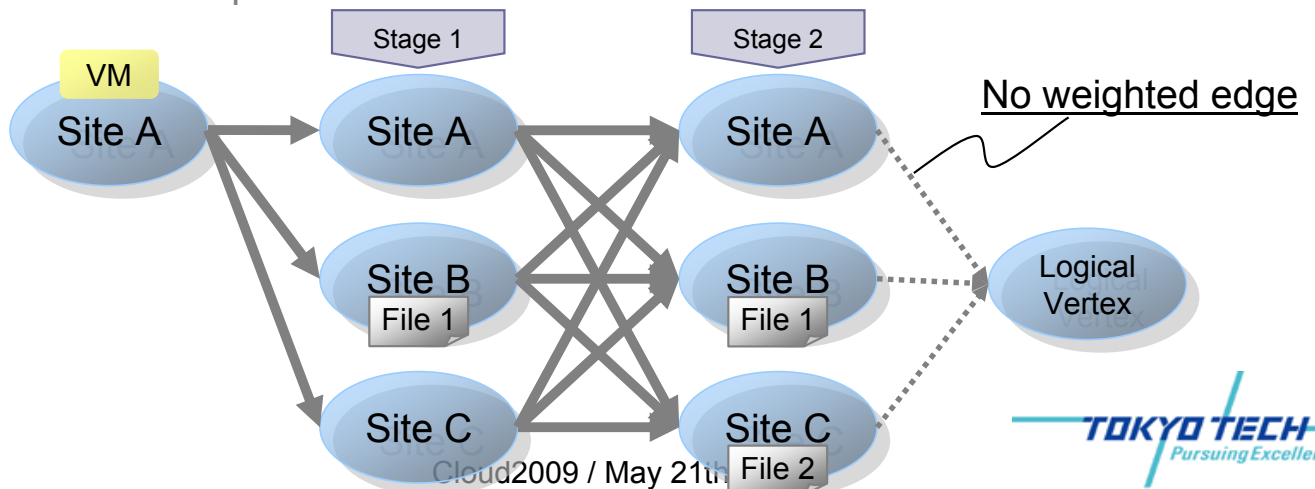
$$P = \begin{pmatrix} 0.2 & 0.8 \\ 1.0 & 0.0 \end{pmatrix} \quad P^2 = \begin{pmatrix} 0.84 & 0.16 \\ 0.2 & 0.8 \end{pmatrix} \quad P^3 = \begin{pmatrix} 0.33 & 0.67 \\ 0.84 & 0.16 \end{pmatrix}$$



How to determine a optimal location for File 1 ?

Search a Shortest Path of the DAG

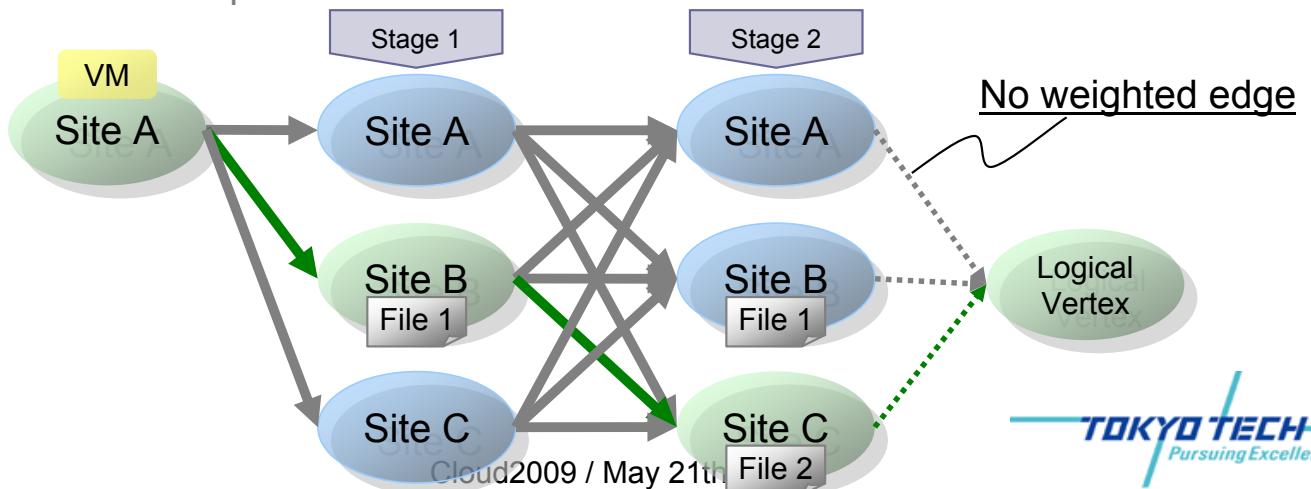
- Adding a logical vertex connected with no weighted edges at the end of DAG
- Solving a shortest path between each ends
 - Vertex weights: Expected file access time
 - Edge weights: VM migration time
- If following path is shortest one
 - Site B is optimal location for File 1 and successive files



How to determine a optimal location for File 1 ?

Search a Shortest Path of the DAG

- Adding a logical vertex connected with no weighted edges at the end of DAG
- Solving a shortest path between each ends
 - Vertex weights: Expected file access time
 - Edge weights: VM migration time
- If following path is shortest one
 - Site B is optimal location for File 1 and successive files



Performance Models

File access time model	$\frac{io_size}{\min(network, local)}$
VM migration time model	$\frac{vm}{network} + c (const)$

io_size : Access File size (MB)

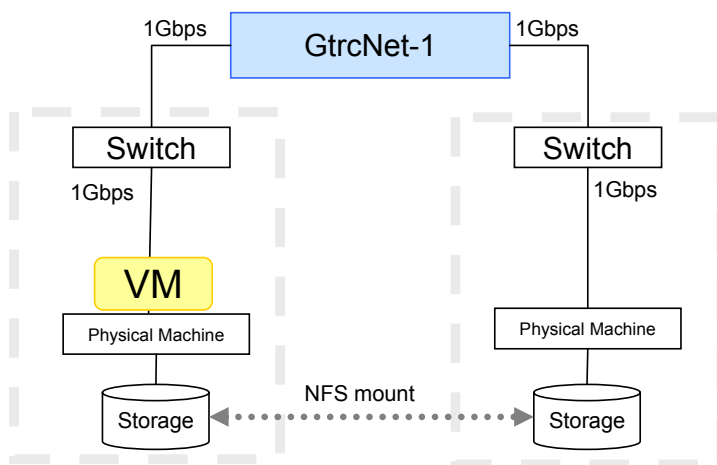
network : Network throughput (MB/s)

local : Local I/O throughput (MB/s)

vm : Allocated VM Memory size (MB)

Experimental Environment for Performance modeling

- Connect 2 machines via network emulator
GtrcNet-1[Kodama et al '04]
 - PrestoIII cluster at Tokyo Tech
- Virtual machine monitor: Xen




- Machines Configurations -

OS	Debian/Linux (kernel: 2.6.18-xen)
CPU	Opteron250 (2.4GHz) * 2
Memory	2GB
NIC	NetXtreme BCM5704
Xen	Xen 3.1.0

Experiments setting for creating Performance Models

- VM Migration Time Model
 - Migrate a VM running an application between two machines
 - Application: BLAST, no application (idle)
 - Network throughputs: 5 – 100 [Mbps]
 - VM memory size: 256, 512, 768, , 1536 [MB]
 - Target VM Migration: Stop-and-Copy Way

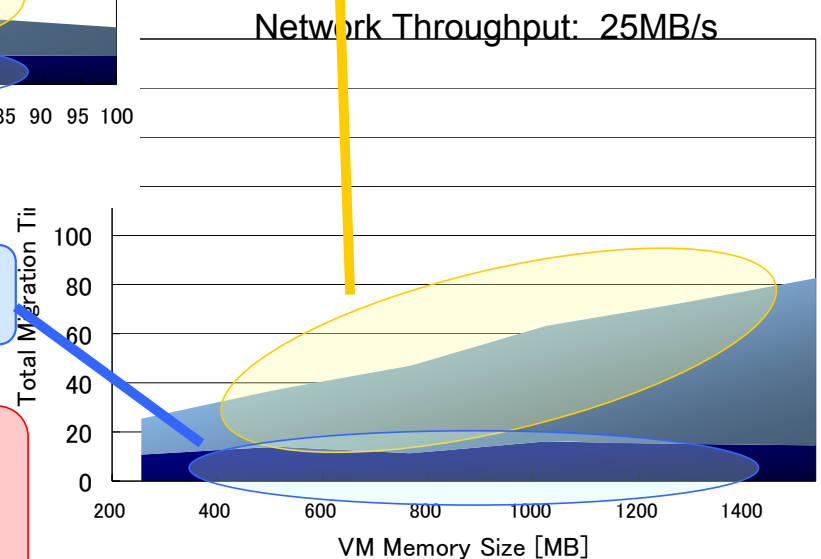
VM Migration time while BLAST exec



1. Memory Image transp time :
 - Inverse proportion to VM memory size
 - Proportion to Network throughput

2. initialization & finalization time is const

3. We got same results in case of migrating VM running no application



Performance Model

File access time model	$\frac{io_size}{\min(network, local)}$
VM migration time model	$\frac{vm}{network} + c (const)$

io_size : Access file size (MB)

$network$: Network throughput (MB/s)

$local$: Local I/O throughput (MB/s)

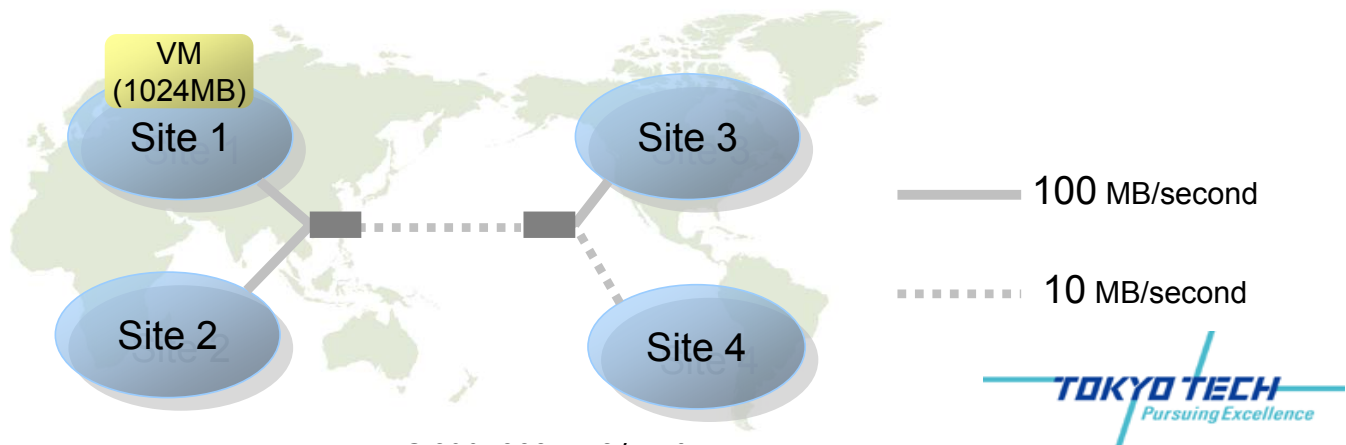
vm : Allocated VM memory size (MB)

Outline

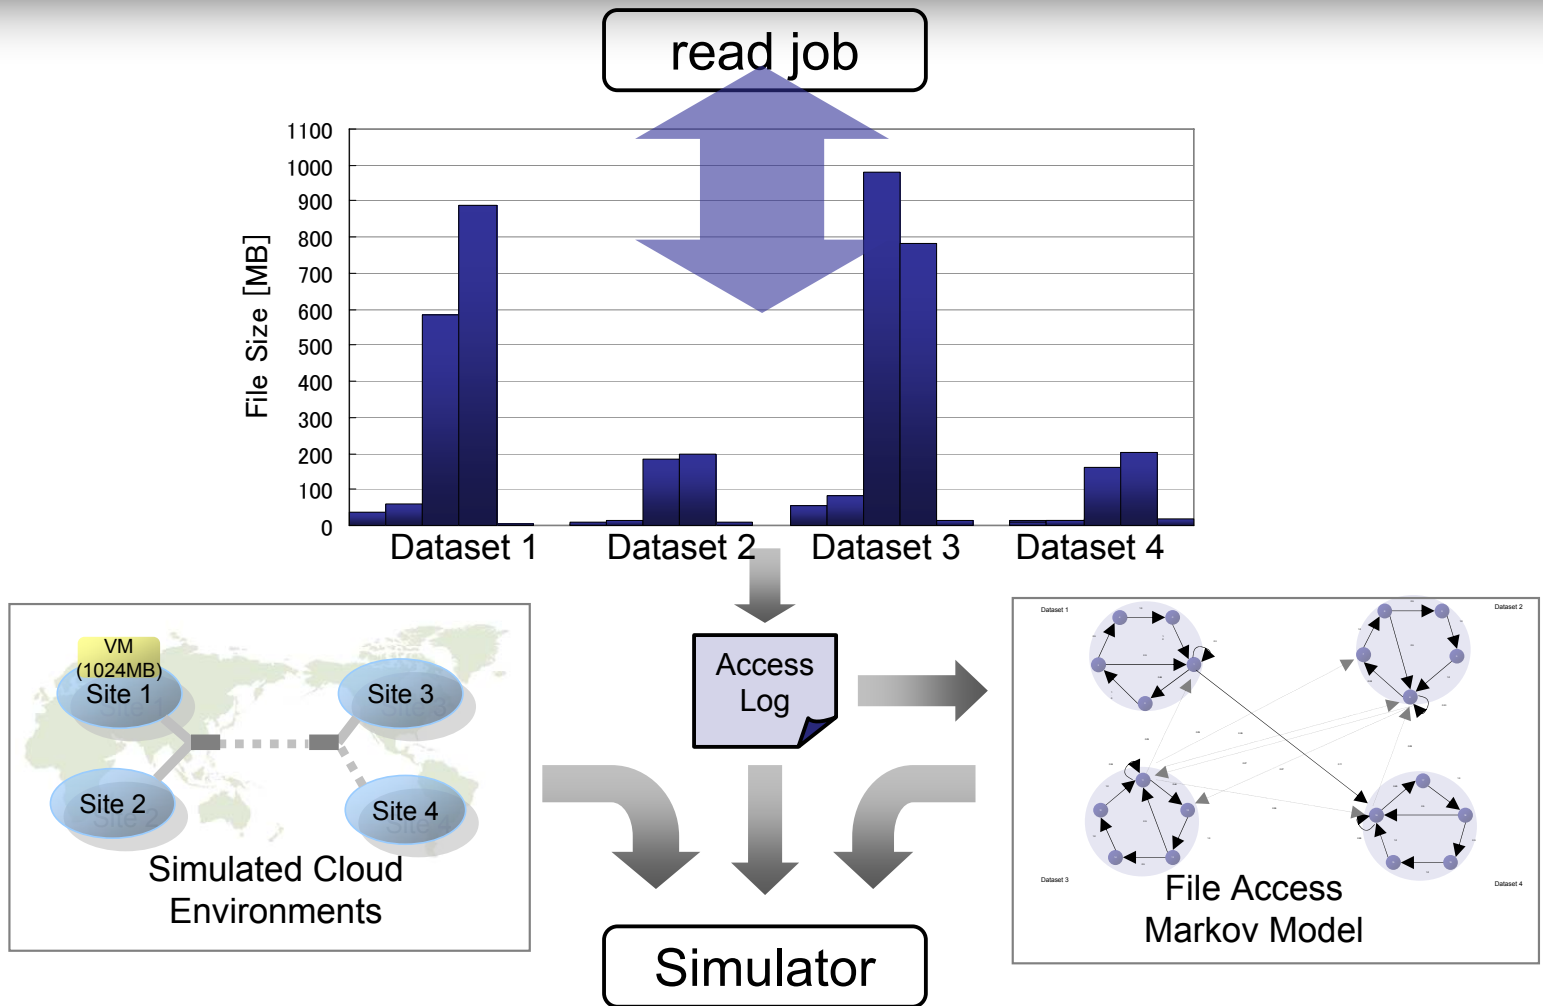
- Introduction
- Target cloud model
- Proposal
 - DAG algorithm
 - Markov model
 - Performance model
- **Evaluation**
- Conclusion

Experiments settings

- Cloud settings
 - Network: 10 or 100[MB/second]
 - Local I/O Throughputs: 60[MB] on each site
- VM settings
 - Memory size: 1024[MB]
 - Initial Location: Site 1




Simulation Method




Markov model of file dependency


Dataset 1




Dataset 2



Dataset 3



Dataset 4



Edge weight (<0.05) is omitted

Experiment targets

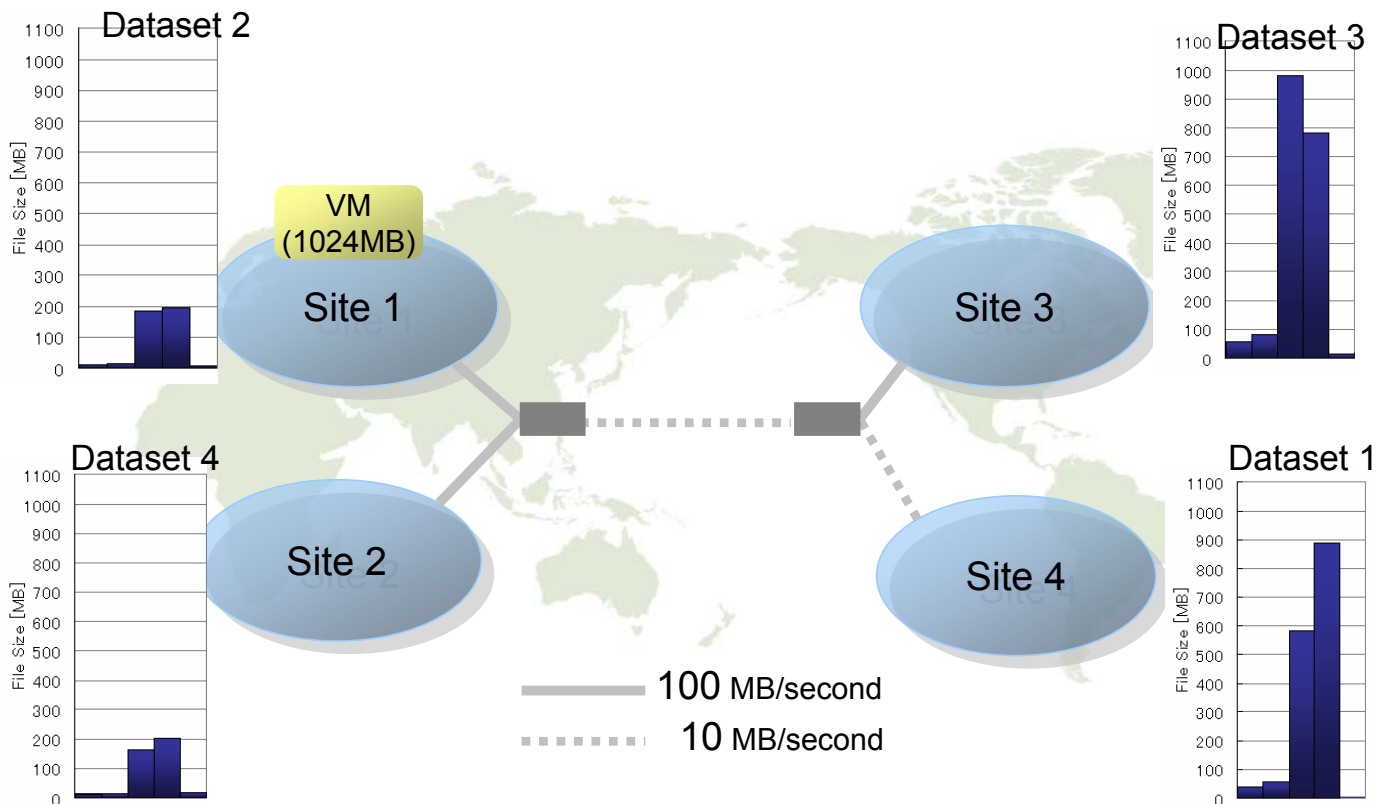
Comparing a Total File Access time with following strategies

- No Migration I/O
 - Always accesses from the initial location (Site 1)
- Migration I/O
 - Always migrates VM onto sites that hold target file
- Proposal
 - Determine the VM migration strategy from our proposed algorithm

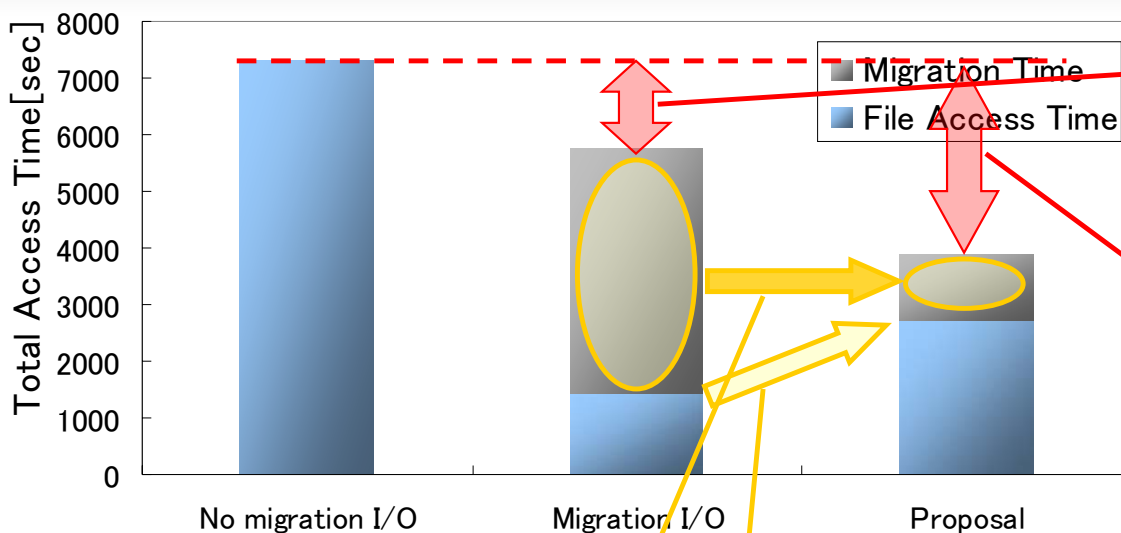
Experiment 1 :

File size & location settings

large size dataset is located far from initial location



Experiment 1: Total file access time

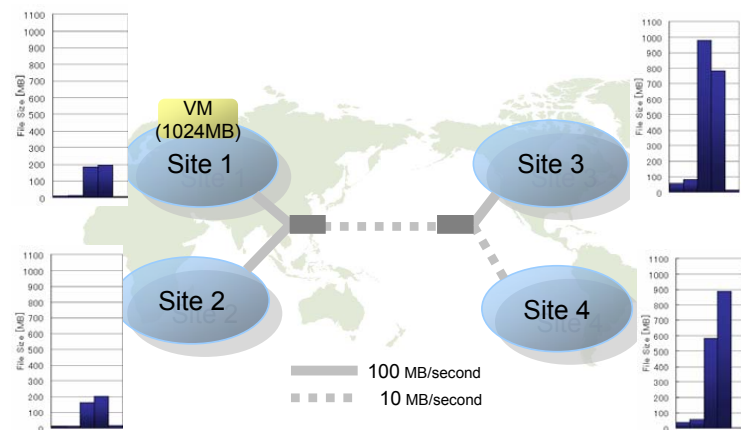


Reducing File access time by 21 %

Reducing File access time by 47 %

Reducing the total file access time more by avoiding unnecessary VM migration in WANs

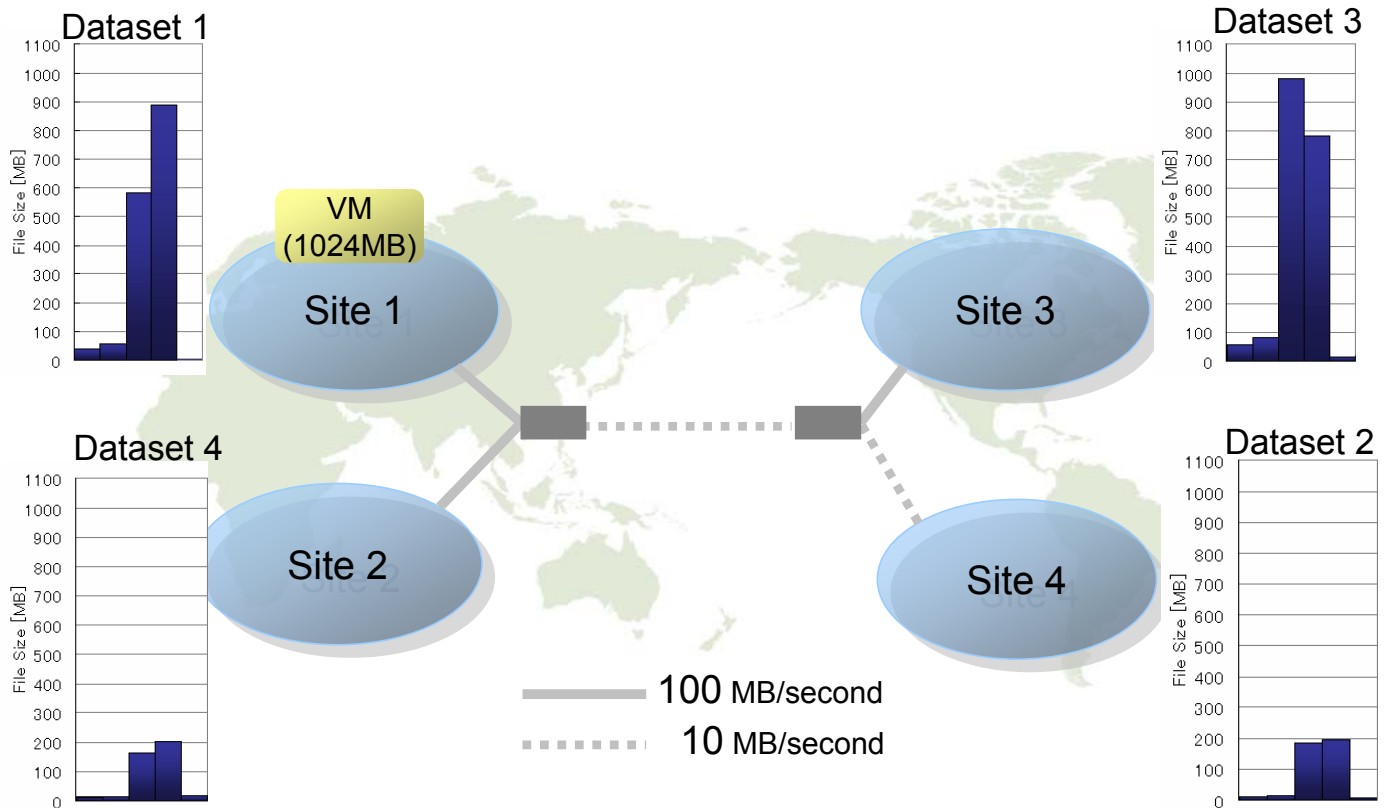
Longer file access time than Migration I/O strategy



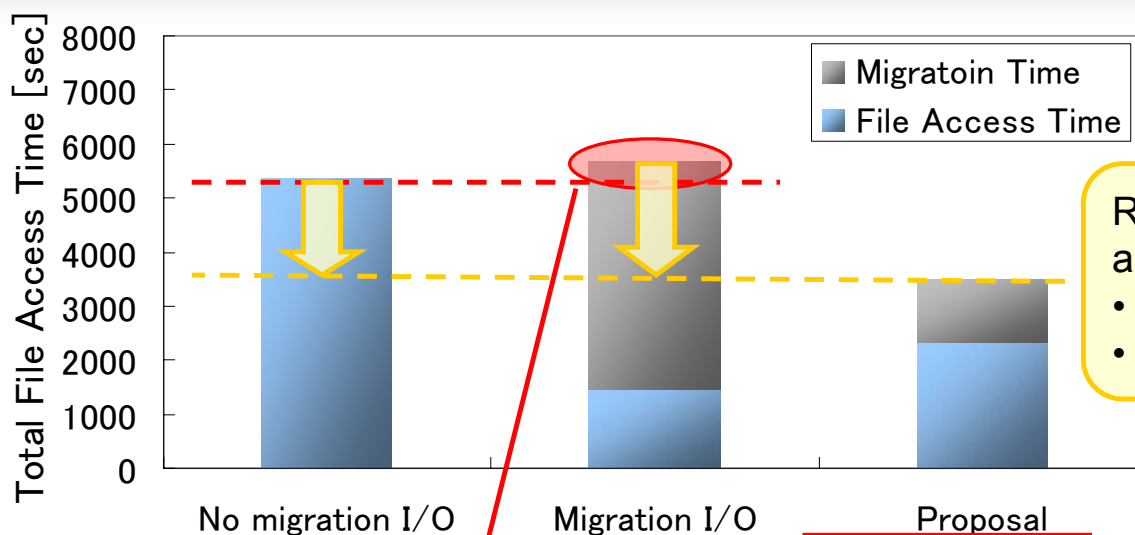
Experiment 2 :

File size & location settings

One of large size dataset is located in initial location



Experiment 2: Total file access time

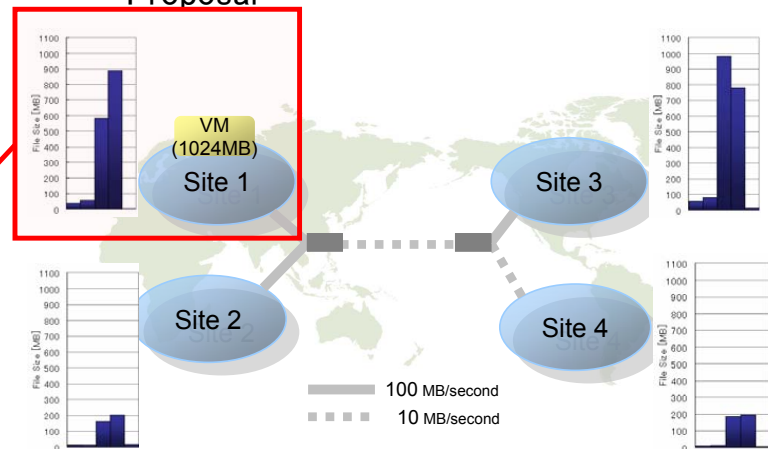


Reducing the total file access time

- No migration I/O: 34%
- Migration I/O: 38 %


Longer than the total access time of the No migration strategy

large size files are located on the same site to the initial location of the VM



Scalability

runtime of the algorithm



Negligibly-small runtime
in this experiment
(4 sites, 5 stages)


- For scalability, we can
 - Set the maximum # of the stages to control the runtime
 - Reuse the results of the shortest path search
 - Solve the shortest path problem previously

Conclusion

- Created Performance model
 - File access time and VM migration time
- Proposed optimizing algorithm for I/O intensive application
 - Representing the access dependency between files as a markov model
 - Determining VM migration strategy
- Achieved higher performance than simple techniques
 - No migration: 38%
 - Always migration: 47%
- Our proposed algorithm is expected to be more effective for applications accessing TB-sized files and larger

Future work

- For the performance model
 - Considering CPU and memory usages for heterogeneous environments
- For the optimizing algorithm
 - Considering other VM placements
 - Load balancing
 - Considering a VM migration algorithm in conjunction with file migration



Thank you,
Any Questions ?

Cloud2009 / May 21th